

GENDER SPEECH RECOGNITION REPORT

2022-07-13

S300733: Atanasio Giuseppe
S296138: Carachino Alessio



**Politecnico
di Torino**

Sommario

Premises	3
Task	3
Dataset	3
Features	4
Histograms	4
Dimensionality Reduction Techniques	7
PCA.....	7
LDA.....	7
Classification	9
What we use... ..	9
Validation	9
Generative models - Multivariate Gaussian Classifier	10
Logistic Regression.....	11
Support Vector Machine	14
Gaussian Mixture Models	16
The wrap up	18
Final picks.....	20
Experimental Results	21
Generative models – Multivariate Gaussian Classifiers	21
Logistic Regression.....	21
Support Vector Machine	23
Gaussian Mixture Model.....	24
Final considerations	25
SVM RBF vs 4-D GMM tied	25
Weighted Logistic Regression vs 4-D GMM tied.....	26

Premises

Task

The goal of the application is to build a model that best fits for the gender classification. We will discuss how they perform for the problem we have chosen, explaining pros and cons.

Dataset

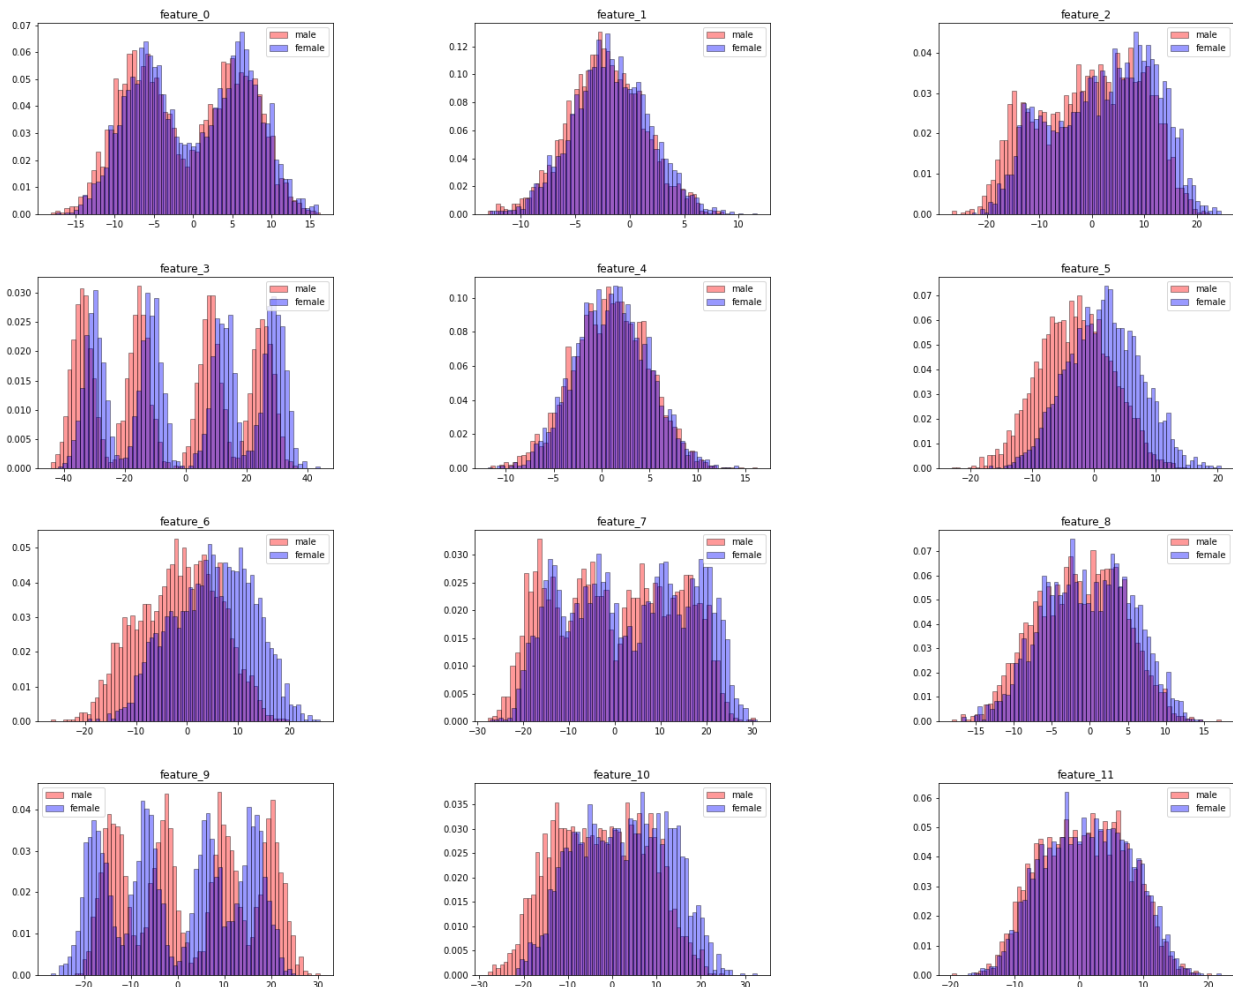
The dataset consists of synthetic speaker embeddings that represent the acoustic characteristics of a spoken utterance. Each row corresponds to a different speaker and contains 12 features followed by the gender label (1 for female, 0 for male). The features do not have any particular interpretation. Speakers belong to four different age groups. The age information, however, is not available.

The training set consists of 3000 samples for each class, whereas the test set contains 2000 samples for each class.


Features

Histograms

Here are histograms for each of the 12 features. We can see that the raw features have an approximated gaussian distribution, but, in the next page, the features' distributions are plotted after they have been pre-processed (Gaussianized).




RAW FEATURES CORRELATION HEATMAP




Feature 3 is highly correlated to the 0, 2, 7 and 9 ones.

This suggests we may benefit from using PCA to map data to less correlated features, but we will check later this assumption.

Here are the histograms of the gaussianized features. As we expected, the gaussianization has not brought so much improvement at all. Instead, gaussians in feature 3 and 9 were better in the previous un-pre-processed dataset.



GAUSSIANIZED FEATURES CORRELATION HEATMAP



Note that, with respect to the non-gaussianized features, the correlation between 7-9 and 3 ones has decreased.

Dimensionality Reduction Techniques

PCA

PCA is the process of computing the principal components and using them to perform a change of basis on the data. This means that it makes it possible to map in another feature space the analysis and even to reduce the number of features of the dataset, reducing for example the training overlay and the risk introduced by the curse of dimensionality.

The following images, which are for demonstration purposes only, show the training set after that PCA has been applied.




Figure 1 - PCA (12 -> 2)

Figure 2 - PCA (12 -> 3)

LDA

While PCA does not provide any guarantee about obtaining discriminant directions, Linear Discriminant Analysis is defined as a criterion of optimality finds a direction that has a large separation between the classes and small spread inside each class.

Also in this case, The following images, which are for demonstration purposes only, show the training set after that LDA has been applied, or PCA followed by LDA, as these two techniques can be used simultaneously (in this strict order).




Figure 3 – PCA (12 -> 2) + LDA (green line is the best direction)




Figure 4 - PCA (12 -> 2) + LDA

(histograms show how points are projected along the direction)




Figure 5 - PCA (12 ->3) + LDA (blue line is the best direction)




Figure 6 - PCA (12 -> 3) + LDA (histograms show how points are projected along the direction)

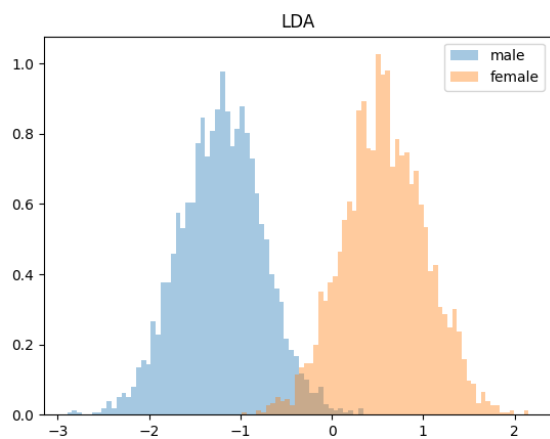


Figure 7 - LDA (12 -> 1)

Classification

What we use...

- Generative models – Linear and Quadratic Classifiers
 - Multivariate Gaussian Classifier (**MVG**)
 - MVG + Diagonal Covariance
 - MVG + Tied Covariance
 - MVG + Diagonal Covariance with Tied Covariance
- Logistic regression
 - Quadratic Logistic Regression
 - Weighted Logistic Regression
- Support Vector Machines (SVM)
 - Primal implementation
 - Kernel SVM with polynomial kernel with degree=4
 - Kernel SVM with RFB
- Gaussian Mixture Models (GMM)
 - Gaussian Mixture Models (**GMM**)
 - GMM + Diagonal Covariance
 - GMM + Tied Covariance
 - GMM + Diagonal Covariance with Tied Covariance

Validation

- To understand which model is most promising, and to assess the effects of using PCA, we have employed K-Fold cross validation. In facts, all of the following results have been obtained with K-Fold Validation with K = 5.
- Inside each cell of the following tables, we have reported the minDCF. We do not care about actDCF in this initial phase.
- 'MinDCF' has been computed with Cfp and Cfn both equal to one, as we do not have any specific requirements regarding the miss-classification costs.

Generative models - Multivariate Gaussian Classifier

EXPECTATIONS

Since histograms have shown that features approximately have a gaussian distribution, it is expected that the Generative Models work well for this dataset. Furthermore, heatmaps showed that correlation is significantly spreaded along the features. Therefore, it is expected that the models based on the Naïve Bayes assumptions will perform badly.

RESULTS

Here are the results of Multivariate Gaussian Classifiers in three different applications (ours has $\pi=0.5$).

	$\pi = 0.5$	$\pi = 0.9$	$\pi = 0.1$	$\pi = 0.5$	$\pi = 0.9$	$\pi = 0.1$	$\pi = 0.5$	$\pi = 0.9$	$\pi = 0.1$
	RAW features			Gaussianization			Z NORM		
	no PCA								
FULL-COV	0.048	0.125	0.128	0.062	0.171	0.181	0.048	0.123	0.126
DIAG-COV	0.563	0.856	0.825	0.541	0.824	0.81	0.565	0.848	0.818
TIED FULL-COV	0.047	0.128	0.118	0.06	0.167	0.18	0.048	0.127	0.125
TIED DIAG-COV	0.564	0.85	0.829	0.538	0.816	0.804	0.566	0.845	0.821
	PCA (m=10)								
FULL-COV	0.047	0.124	0.14	0.071	0.204	0.206	0.112	0.263	0.304
DIAG-COV	0.067	0.156	0.162	0.085	0.223	0.228	0.119	0.275	0.294
TIED FULL-COV	0.047	0.125	0.13	0.071	0.206	0.199	0.111	0.259	0.298
TIED DIAG-COV	0.063	0.149	0.153	0.083	0.225	0.227	0.118	0.272	0.293
	PCA (m=9)								
FULL-COV	0.047	0.125	0.139	0.091	0.238	0.242	0.158	0.376	0.403
DIAG-COV	0.065	0.153	0.164	0.095	0.258	0.26	0.161	0.377	0.417
TIED FULL-COV	0.047	0.123	0.131	0.09	0.233	0.236	0.155	0.367	0.398
TIED DIAG-COV	0.062	0.146	0.153	0.096	0.261	0.26	0.161	0.376	0.415

CONSIDERATIONS

As it was expected, PCA brings good results even with 9 features and improves models based on Naïve assumption, but they aren't remarkable with respect to the full and tied RAW features. So, for the discriminative approaches, it's not expected that PCA works well; however, we reported the results anyway. Tied Models works well, and this suggests that models which exploits linear separation rules should work well. This will be confirmed later.

Lastly, Z-Normalized features does not bring remarkable results.

We can notice that the ones with diagonal covariance matrices perform worse than full covariance ones, and this is due to the highly correlation between features, as we have assumed above with the expectations.

Logistic Regression


EXPECTATIONS

Since 'TIED FULL-COV' introduces linear separation rules and have good results, it is expected that Logistic Regression will perform well. Furthermore, we do not expect that with 'Gaussianization' or 'Z-score Normalization' we will obtain significantly better results as Logistic Regression does not require specific assumptions on the data distribution.

RESULTS

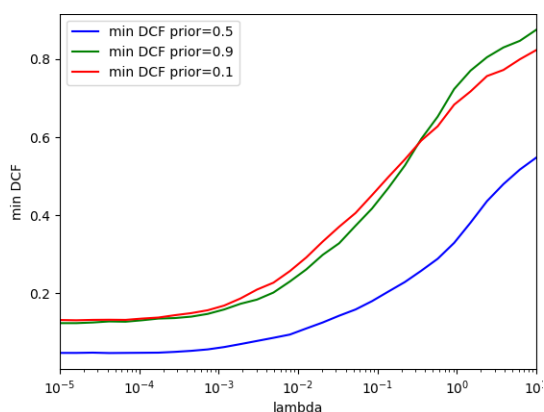
The plots show how minDCF is affected by different values of λ . They are exploited to calibrate λ , which is the regularization term.

LR RAW FEATURES – Results




	$\pi = 0.5$	$\pi = 0.9$	$\pi = 0.1$
NO PCA			
LR, LAMBDA=1E-6	0.048	0.126	0.131
LR, LAMBDA=1E-4	0.048	0.126	0.131
LR, LAMBDA=1E-2	0.048	0.126	0.132
LR, LAMBDA=1.0	0.068	0.157	0.176
PCA m=10			
LR, LAMBDA=1E-6	0.048	0.122	0.139
LR, LAMBDA=1E-4	0.048	0.122	0.139
LR, LAMBDA=1E-2	0.048	0.12	0.136
LR, LAMBDA=1.0	0.07	0.159	0.174
PCA m=9			
LR, LAMBDA=1E-6	0.048	0.124	0.137
LR, LAMBDA=1E-4	0.048	0.124	0.137
LR, LAMBDA=1E-2	0.048	0.122	0.134
LR, LAMBDA=1.0	0.07	0.158	0.174

LR Z-NORMALIZED FEATURES – Results



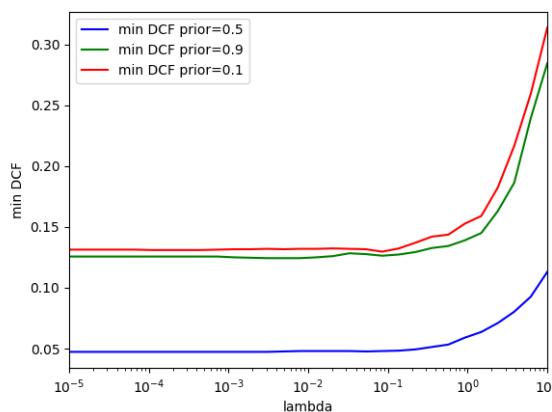
	$\Pi = 0.5$	$\Pi = 0.9$	$\Pi = 0.1$
NO PCA			
LR, LAMBDA=1E-6	0.048	0.126	0.131
LR, LAMBDA=1E-4	0.047	0.13	0.134
LR, LAMBDA=1E-2	0.101	0.239	0.273
PCA m=10			
LR, LAMBDA=1E-6	0.112	0.26	0.296
LR, LAMBDA=1E-4	0.114	0.261	0.301
LR, LAMBDA=1E-2	0.131	0.317	0.342
PCA m=9			
LR, LAMBDA=1E-6	0.158	0.359	0.405
LR, LAMBDA=1E-4	0.158	0.358	0.404
LR, LAMBDA=1E-2	0.164	0.397	0.432

LR GAUSSIANIZED FEATURES – Results



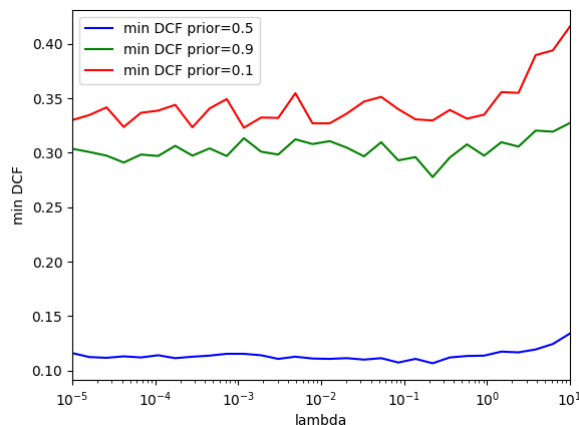
	$\Pi = 0.5$	$\Pi = 0.9$	$\Pi = 0.1$
NO PCA			
LR, LAMBDA=1E-6	0.056	0.159	0.163
LR, LAMBDA=1E-4	0.056	0.158	0.165
LR, LAMBDA=1E-2	0.08	0.203	0.226
PCA m=10			
LR, LAMBDA=1E-6	0.068	0.198	0.194
LR, LAMBDA=1E-4	0.068	0.202	0.192
LR, LAMBDA=1E-2	0.088	0.23	0.234
PCA m=9			
LR, LAMBDA=1E-6	0.089	0.237	0.232
LR, LAMBDA=1E-4	0.089	0.238	0.233
LR, LAMBDA=1E-2	0.099	0.255	0.276

WEIGHTED LR RAW FEATURES – Results



	$\Pi = 0.5$	$\Pi = 0.9$	$\Pi = 0.1$
NO PCA			
WLR, LAMBDA=1E-4	0.047	0.126	0.131
WLR, LAMBDA=1E-2	0.048	0.125	0.132
WLR, LAMBDA=1E-1	0.048	0.127	0.13
PCA m=10			
WLR, LAMBDA=1E-4	0.048	0.122	0.139
WLR, LAMBDA=1E-2	0.048	0.12	0.137
WLR, LAMBDA=1E-1	0.048	0.125	0.134
PCA m=9			
WLR, LAMBDA=1E-4	0.048	0.124	0.137
WLR, LAMBDA=1E-2	0.048	0.122	0.136
WLR, LAMBDA=1E-1	0.048	0.124	0.135

QUADRATIC LR RAW FEATURES – Results



	$\Pi = 0.5$	$\Pi = 0.9$	$\Pi = 0.1$
NO PCA			
QLR, LAMBDA=1E-4	0.117	0.299	0.322
QLR, LAMBDA=1E-2	0.112	0.309	0.351
QLR, LAMBDA=1E-1	0.107	0.302	0.345
PCA m=10			
QLR, LAMBDA=1E-4	0.048	0.122	0.139
QLR, LAMBDA=1E-2	0.048	0.12	0.136
QLR, LAMBDA=1E-1	0.049	0.13	0.133
PCA M=9			
QLR, LAMBDA=1E-4	0.048	0.124	0.137
QLR, LAMBDA=1E-2	0.048	0.122	0.134
QLR, LAMBDA=1E-1	0.049	0.127	0.134

CONSIDERATIONS

The calibration of the hyperparameter λ shows that regularization is not required in none of the above cases. However, we decided to show the minDCF value for different λ values. Results for $\lambda=0$ are not reported since they are similar to $\lambda=1e-6$ or $\lambda=1e-4$. PCA does not bring remarkable improvements, except for the Quadratic Logistic Regression.

The best performing model of Log-Reg family on the validation set is the Weighted Logistic Regression with no regularization (or λ close to zero).

Up to now, the top 2 performing models are Weighted Log-Reg and MVG with same covariance matrixes (Tied).

Support Vector Machine

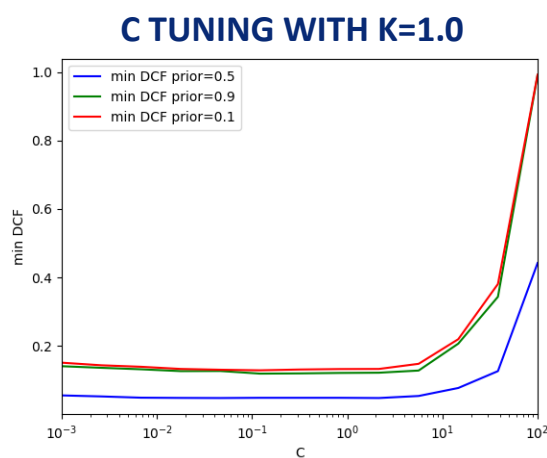
EXPECTATIONS

Up to now, model which exploit linear separation rules have performed very well. So it is expected that Linear SVM (the primal implementation) will provide good results as well.

RESULTS

The plots show how minDCF is affected by different values of the hyper-parameter of C. They are exploited to tune C.

Primal, RAW FEATURES Implementation – Results



$\pi = 0.5$	$K = 0.1$	$K = 1.0$	$K = 10.0$
SVM, C=0.01	0.06	0.049	0.048
SVM, C=0.1	0.057	0.048	0.048
SVM, C=1.0	0.051	0.048	0.047
SVM, C=10.0	0.093	0.096	0.065

$\pi = 0.1$	$K = 0.1$	$K = 1.0$	$K = 10.0$
SVM, C=0.01	0.163	0.135	0.132
SVM, C=0.1	0.161	0.13	0.13
SVM, C=1.0	0.14	0.133	0.138
SVM, C=10.0	0.252	0.232	0.165

$\pi = 0.9$	$K = 0.1$	$K = 1.0$	$K = 10.0$
SVM, C=0.01	0.151	0.13	0.124
SVM, C=0.1	0.146	0.12	0.118
SVM, C=1.0	0.129	0.119	0.127
SVM, C=10.0	0.291	0.291	0.157

Dual + Polynomial, RAW FEATURES Implementation – Results

C = 1.0, D=4, CON=1000	K = 1.0	K = 10.0
SVM, $\Pi=0.5$	0.11	0.166
SVM, $\Pi=0.1$	0.334	0.543
SVM, $\Pi=0.9$	0.322	0.44

Dual + RBF, RAW FEATURES Implementation – Results

GAMMA=0.001 C=1.0	K = 0.1	K = 1.0	K = 10.0
SVM, $\Pi = 0.5$,	0.039	0.039	0.04
SVM, $\Pi = 0.1$	0.122	0.123	0.12
SVM, $\Pi = 0.9$	0.12	0.12	0.12

CONSIDERATIONS

We can compare linear models in terms of min DCF:

	$\Pi = 0.5$	$\Pi = 0.9$	$\Pi = 0.1$
MVG (TIED FULL-COV)	0.047	0.128	0.118
LOG REG ($\Lambda = 10^{-6}$, $\Pi T = 0.5$)	0.048	0.126	0.131
WLOG REG ($\Lambda = 10^{-4}$, $\Pi T = 0.5$)	0.047	0.126	0.131
SVM, C=1.0	0.048	0.119	0.133

Linear SVM performs similarly to other linear approaches, as expected.

Since Linear models perform similarly to non-linear models, we decided to report the results for the two SVM non-linear formulations

Gaussian Mixture Models

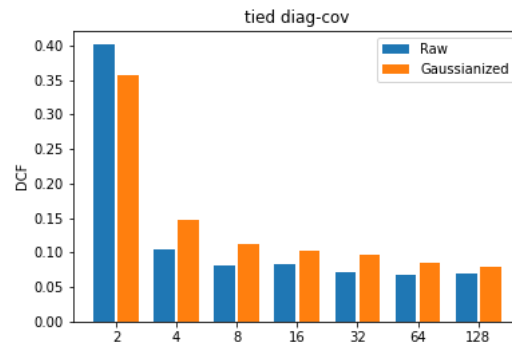
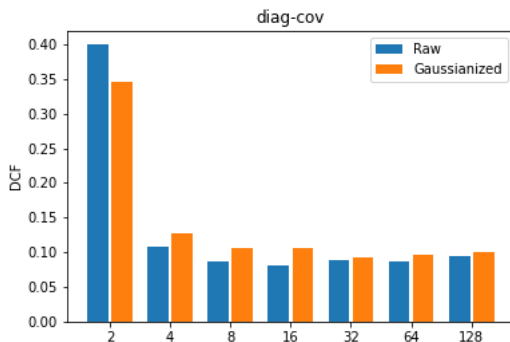
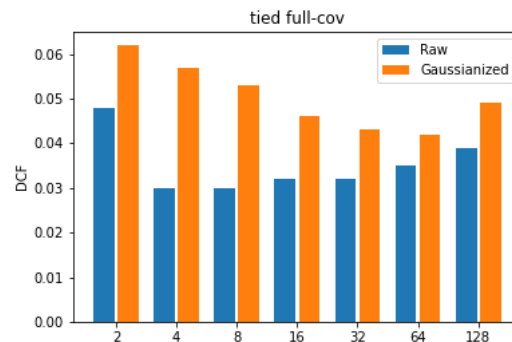

EXPECTATIONS

GMMs can approximate generic distributions, so we expect to obtain better results than with the Gaussian model. We are going to train GMM up to 128 components for full and diagonal covariance, and tied/non tied.

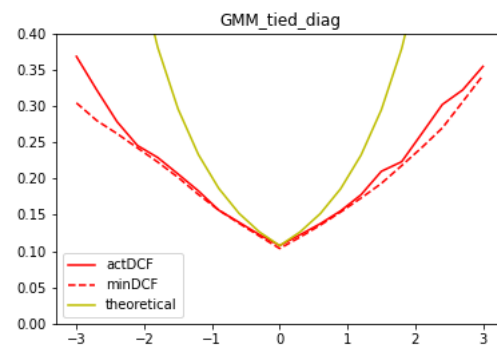
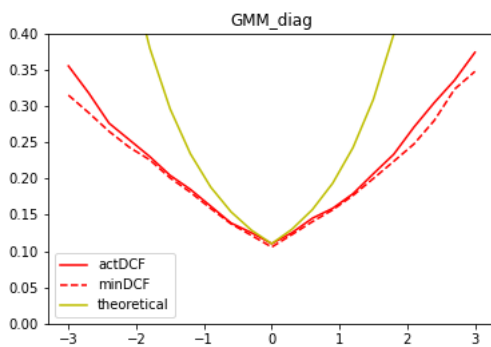

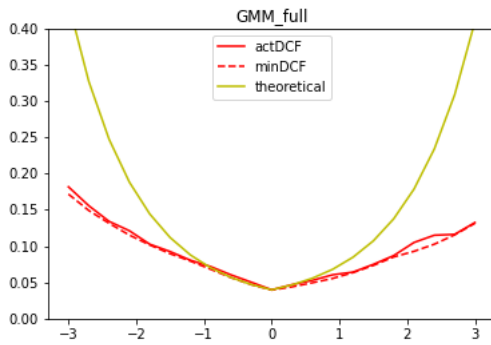
Please note: raw and z-normalization retrieve the same results.

RESULTS

COMPONENTS	2	4	8	16	32	64	128
	Z-Normalization						
FULL-COV	0.043	0.033	0.032	0.044	0.057	0.087	0.151
DIAG-COV	0.4	0.108	0.087	0.081	0.088	0.086	0.095
TIED FULL-COV	0.048	0.03	0.03	0.032	0.035	0.035	0.039
TIED DIAG-COV	0.402	0.105	0.082	0.083	0.072	0.068	0.069
	Gaussianized						
FULL-COV	0.053	0.047	0.048	0.055	0.064	0.089	0.106
DIAG-COV	0.346	0.127	0.106	0.107	0.093	0.096	0.101
TIED FULL-COV	0.062	0.057	0.053	0.046	0.043	0.042	0.049
TIED DIAG-COV	0.358	0.148	0.113	0.103	0.098	0.086	0.08



4-D GMM



CONSIDERATIONS

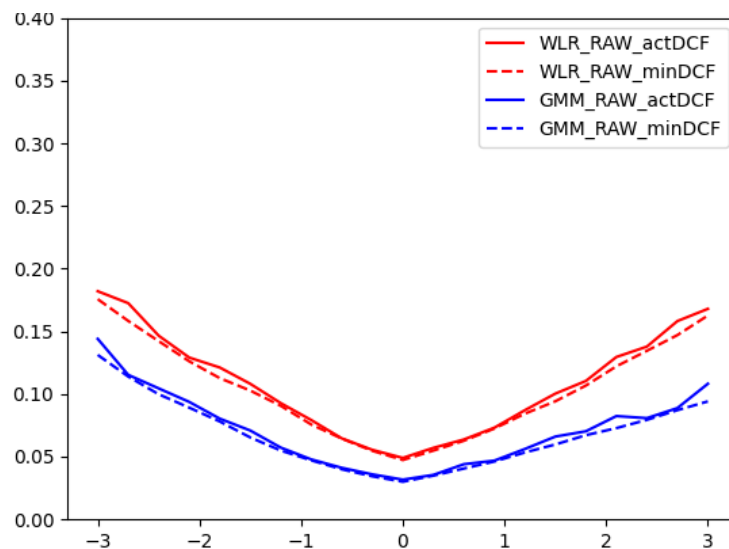
We can check that the best result that it could be obtained is for full-covariance tied GMM, especially the 4-D one.

Taking a zoom for the 4-D models, we can analyse bayes error plots: the best is again the full-covariance tied.


The wrap up

We still should compare the models. It's time to exploit ROC curves and Bayes Error Plots.

Firstly, here is the comparison (on RAW features) between the GMM with 4 components and Tied Full Covariance Matrix and Weighted Logistic Regression.




It is clear that GMM outperforms Weighted Logistic Regression. However, calibration is not required as the actDCF curve is already quite close to the minDCF curve.

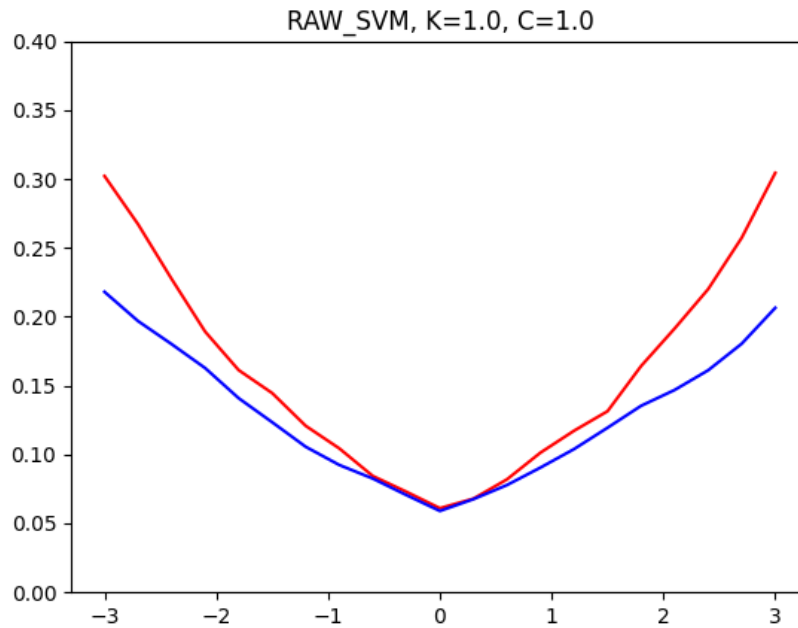


The ROC Curve shows the better performance on the validation set of the GMM too.

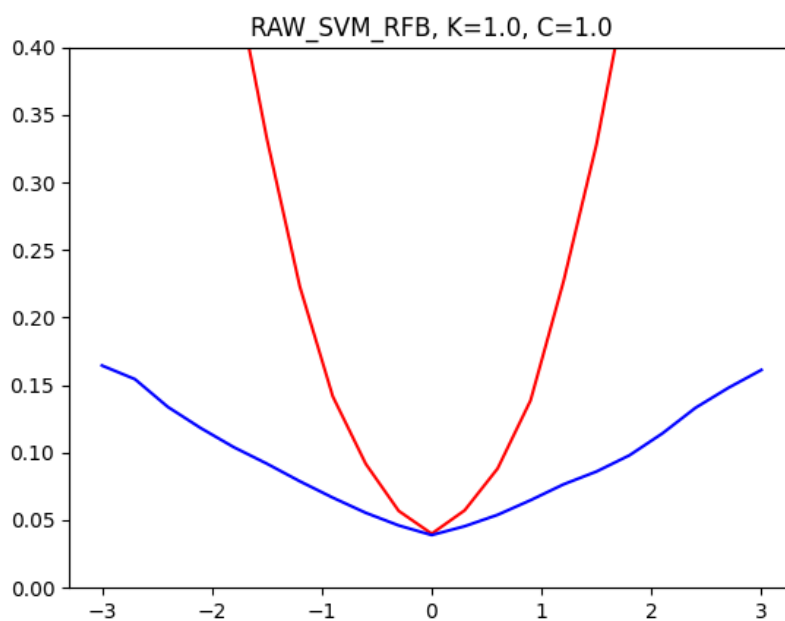
Now, let's turn the attention on the comparison (RAW features in this case too) between MGC with Tied Full Covariance and the same GMM as above.



Where the exact same considerations of the previous comparison are still valid. For what regards SVM, the following image shows the Bayes Error Plots for the case of the linear SVM:



The above results are good but not better than the previously compared models. While the next one, shows the Bayes Error Plot for the RFB Kernel case.



In this case, scores are highly uncalibrated. They can be improved by basically applying Logistic Regression on the scores returned by the model.

Final picks

We expect that the **4D-GMM Tied Full Covariance Model** outperforms the others.

Experimental Results

We now turn the attention to the evaluation part. It's useful to choose the minimum DCF to evaluate because it's an optimistic estimation for the actual DCF in the hypothetical case in which it could be chosen the optimal threshold for the evaluation set.

Since validation has performed with 5-fold cross-validation, the evaluation part will be executed using the whole dataset (all training data for training, all test data for test).

Generative models – Multivariate Gaussian Classifiers

Here are reported the evaluation results for MVG.

	$\Pi = 0.5$	$\Pi = 0.9$	$\Pi = 0.1$	$\Pi = 0.5$	$\Pi = 0.9$	$\Pi = 0.1$
	RAW features			Gaussianization		
				no PCA		
FULL-COV	0.053	0.138	0.134	0.073	0.182	0.202
DIAG-COV	0.57	0.882	0.81	0.547	0.846	0.792
TIED FULL-COV	0.05	0.135	0.132	0.069	0.177	0.184
TIED DIAG-COV	0.57	0.88	0.808	0.545	0.846	0.793
				PCA (m=10)		
FULL-COV	0.052	0.139	0.136	0.08	0.206	0.203
DIAG-COV	0.068	0.18	0.198	0.088	0.229	0.243
TIED FULL-COV	0.054	0.134	0.138	0.078	0.205	0.199
TIED DIAG-COV	0.068	0.181	0.196	0.085	0.225	0.234
				PCA (m=9)		
FULL-COV	0.054	0.14	0.144	0.098	0.23	0.255
DIAG-COV	0.068	0.18	0.204	0.102	0.248	0.29
TIED FULL-COV	0.054	0.136	0.142	0.098	0.232	0.258
TIED DIAG-COV	0.068	0.18	0.2	0.104	0.258	0.286

The results are consistent with those obtained in the validation part. We can confirm that the MVG model with tied covariance is the best option to choose, and this also confirm that linear separation rules are the best choose.

PCA isn't outperforming with respect to the previous, even if it improves the MVG models with diagonal covariance.

Gaussianization worsen the data.

Logistic Regression

We now move our attention to the Logistic Regression.

	$\Pi = 0.5$	$\Pi = 0.9$	$\Pi = 0.1$
LR RAW LAMBDA=1E-6	0.053	0.133	0.135
LR GAUSS. LAMBDA=1E-6	0.059	0.165	0.162
LR ZNORM. LAMBDA=1E-6	0.052	0.135	0.133
WLR RAW, LAMBDA=1E-4	0.053	0.133	0.135
QLR RAW, LAMBDA=1E-4	0.114	0.285	0.353

Even these results are consistent with our expectations.



The results are quite similar to the MVG tied, and it can be confirmed, also in this case, that linear separation rules are the best road to follow.

Although we do not expect significant deviations, we can also verify whether the chosen λ provides close to optimal results.



VALIDATION RESULTS

EVALUATION RESULTS



**LR
RAW**



**WLR
RAW**



**QLR
RAW**



Support Vector Machine

We report the results for all types of SVM, in particular for the hyperparameters $C=1.0$ and $K=1.0$

	$\Pi = 0.5$	$\Pi = 0.9$	$\Pi = 0.1$
SVM PRIMAL	0.051	0.129	0.138
SVM DUAL, POLYNOMIAL D=4	0.156	0.392	0.486
SVM DUAL, RFB GAMMA=1E-3	0.044	0.112	0.133





We reported the comparison on the evaluation set of SVM in the next pages.

Gaussian Mixture Model

We have chosen 4 components results to analyze with other applications.

Please note: raw and z-normalization retrieve the same results.

	$\pi = 0.5$	$\pi = 0.1$	$\pi = 0.9$
	RAW		
FULL-COV	0.041	0.114	0.119
DIAG-COV	0.112	0.258	0.29
TIED FULL-COV	0.03	0.08	0.092
TIED DIAG-COV	0.111	0.256	0.295
	Gaussianization		
FULL-COV	0.05	0.149	0.147
DIAG-COV	0.132	0.328	0.357
TIED FULL-COV	0.06	0.16	0.176
TIED DIAG-COV	0.161	0.411	0.399



The outcomes are consistent with those in the validation part.

We can confirm that linear models outperform, especially for the 4-D GMM tied full covariance case.

It is remarkable that full-covariance models over the 8-D tends to have more problems of overfitting than the diagonal ones.

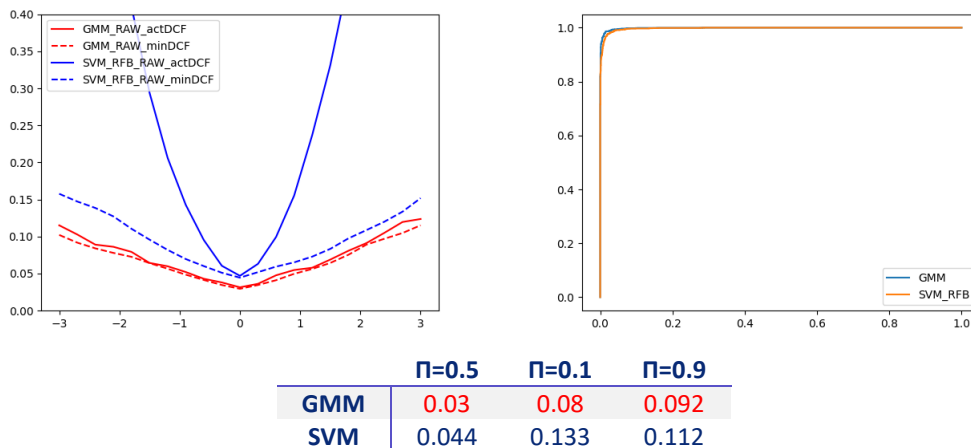
Final considerations

As we expected, the evaluation confirms every result obtained during the validation phase. We can conclude that our results are reliable.

4-D GMM Tied outperforms with respect to the other models.

SVM RBF vs 4-D GMM tied


We now turn our attention to the comparison between SVM RBF and GMM tied. We have chosen 4-G Tied full-cov and SVM dual RBF ($C=1.0$, $K=1.0$). From now on we are going to call them simply SVM and GMM. Following: bayes error (left), ROC (right), comparison (bottom).



As we can see, GMM performs better than SVM. We could calibrate the scores of the latter exploiting logistic regression and check if it performs better, even though GMM remains the best model with the lowest minDCF and actDCF curves. We can conclude that calibrating SVM isn't useful to do at all.

Weighted Logistic Regression vs 4-D GMM tied

We now turn our attention to the comparison between **Weighted Logistic Regression (lambda=1e-4)** and **4-D GMM tied**. For simplicity we will call them WLR and GMM. Following: bayes error (left), ROC (right), comparison (bottom).



	$\pi=0.5$	$\pi=0.1$	$\pi=0.9$
GMM	0.03	0.08	0.092
WLR	0.053	0.135	0.133

Even though WLR has a lower minDCF than the above SVM, the actDCF curve suggests that it performs better than the above SVM model, but not enough to bring us to **choose 4-D GMM tied full-cov as the best model for our application.**